

DSToolKit - Text2SQL and Azure AI Search

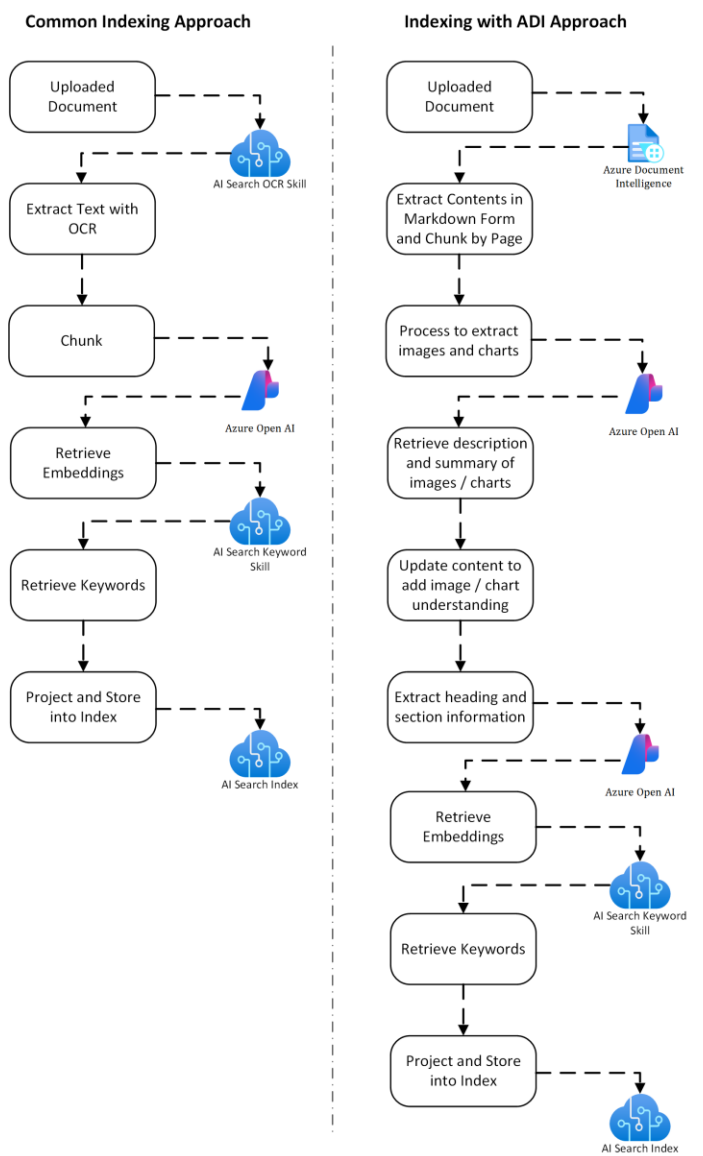


Document Cracking With ADI

Unstructured Data Search

Industry Solutions | Data & AI

Unstructured Search - Document Cracking Approach



Unstructured Search Index Set Up

rag-documents-index ...


	Id	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
	Title	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾
🔑	ChunkId	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/> Keywo... ▾
	Chunk	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾
	Sections	StringCollection	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾
⚡	ChunkEmbedding	SingleCollection	<input type="checkbox"/>				<input checked="" type="checkbox"/> 1536
	Keywords	StringCollection	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾
	SourceUri	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾
>	Figures	ComplexTypeColle...					
	DateLastModified	DateTimeOffset	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	
	PageNumber	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/> Standa... ▾

Document Cracking with Figure Analysis

Model	Ctx Size	Python	C++	Rust	Java	TypeScript	Average
gpt-4o-2024-05-13	128k	95	80	85	96	97	90.6
gemin-1.5-flash-latest	1000k	93	79	87	94	97	90
Phi-3.5-MoE	80	128k	89	74	81	88	95
Phi-3.5-Mini	128k	86	67	73	76	82	77
Llama-3.1-8B-Instruct	128k	80	65	73	76	76	71
Mixtral-8x7B-Instruct-v0.1	32k	66	65	64	71	74	68
Mixtral-8x22B-Instruct-v0.1	64k	60	67	74	83	55	67.8

Table 1: Comparison results on RepoQA benchmark.

such as Arabic, Chinese, Russian, Ukrainian, and Vietnamese, with average MMLU-multilingual scores of 55.4 and 47.3, respectively. Due to its larger model capacity, **phi-3.5-MoE** achieves a significantly higher average score of 69.9, outperforming **phi-3.5-mini**.

Figure 4: Comparison of **phi-3-mini**, **phi-3.5-mini** and **phi-3.5-MoE** on MMLU-Multilingual tasks

We evaluate the **phi-3.5-mini** and **phi-3.5-MoE** models on two long-context understanding tasks: RULER [HSK*24] and RepoQA [LTD*24]. As shown in Tables 1 and 2, both **phi-3.5-MoE** and **phi-3.5-mini** outperform other open-source models with larger sizes, such as Llama-3.1-8B, Mixtral-8x7B, and Mixtral-8x22B, on the RepoQA task, and achieve comparable performance to Llama-3.1-8B on the RULER task. However, we observe a significant performance drop when testing the 128K context window on the RULER task. We suspect this is due to the lack of high-quality long-context data in mid-training, an issue we plan to address in the next version of the model release.

In the table 3, we present a detailed evaluation of the **phi-3.5-mini** and **phi-3.5-MoE** models compared with recent SoTA pretrained language models, such as GPT-4o-mini, Gemini-1.5 Flash, and open-source models like Llama-3.1-8B and the Mistral models. The results show that **phi-3.5-mini** achieves performance comparable to much larger models like Mistral-Nemo-12B and Llama-3.1-8B, while **phi-3.5-MoE** significantly outperforms other open-source models, offers performance comparable to Gemini-1.5 Flash, and achieves above 90% of the average performance of GPT-4o-mini across various language benchmarks.

[illegible]

Figure Understanding with gpt4o to Update Markdown

Chunk Comparison (Old vs New Indexing Approach)

"Chunk": "as mapping a query and a set of key-value pairs to an output, where the query, keys, values, and output are all vectors. The output is computed as a weighted sum of the values, where the weight assigned to each value is computed by a compatibility function of the query with the corresponding key. 2.1 Scaled Dot-Product Attention We call our particular attention "Scaled Dot-Product Attention" (Figure 2). The input consists of queries and keys of dimension d_k , and values of dimension d_v . We compute the dot products of the query with all keys, divide each by $\sqrt{d_k}$, and apply a softmax function to obtain the weights on the values. In practice, we compute the attention function on a set of queries simultaneously, packed together into a matrix Q . The keys and values are also packed together into matrices K and V . We compute the matrix of outputs as: $\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$. The two most commonly used attention functions are additive attention [2], and dot-product (multi-head) attention. Dot-product attention is identical to our algorithm, except for the scaling factor of $1/\sqrt{d_k}$. Additive attention computes the compatibility function using a feed-forward network with a single hidden layer. While the two are similar in theoretical complexity, dot-product attention is much faster and more space-efficient in practice, since it can be implemented using highly optimized matrix multiplication code. While for small values of d_k the two mechanisms perform similarly, additive attention outperforms dot product attention without scaling for larger values of d_k [3]. We suspect that for large values of d_k , the dot products grow large in magnitude, pushing the softmax function into regions where it has extremely small gradients."

"Chunk": "Scaled Dot-Product Attention FigureId="4.1" FigureContent="This image depicts a component of the Transformer model architecture, specifically the attention mechanism. Here is a detailed analysis of the image: Analysis: Modules Displayed: MatMul: The first module shown is a matrix multiplication operation. SoftMax: This operation normalizes the matrix values into probabilities. Mask (optional): An optional masking operation, commonly used to mask certain parts of the input matrices, for instance, to prevent attending to future positions in sequence-to-sequence models. Scale: This operation scales the matrix values, typically by a constant factor. MatMul: There is another matrix multiplication operation at the bottom of the stack. Inputs: Q (Query): The input query vector. K (Key): The input key vector. V (Value): The input value vector. Process Flow: The process starts with scaling the dot products of the query and key matrices (MatMul). The resultant matrix is then masked if necessary. The scaled scores are processed through the softmax function to obtain the attention weights. Finally, these weights are used in the second MatMul operation with the value matrix to get the attention output. Functional Understanding: MatMul (Q, K^T): This computes the dot product of the query and key matrices, providing a raw score for attention. Scale: The scaling (often by the square root of the key vector dimension) helps in maintaining stable gradients. Mask (optional): This optional step masks certain positions, typically for preventing attendance to future tokens in tasks like sequence prediction. SoftMax: Normalizes the scores to probabilities, emphasizing the most relevant parts of the input. MatMul (output, V): This combines the attention weights with the values to get the final context-aware representation. Technical Implications: Efficiency: The use of these modular operations allows efficient parallel processing, key to the speed advantage of transformers. Attention Mechanism: The detailed attention mechanism allows the model to weigh input components differentially, crucial for capturing dependencies in data sequences. Scalability: The components' modular nature makes the architecture easily scalable to larger datasets and complex tasks. Conclusion: The image represents the core structure of the attention mechanism from the Transformer model, demonstrating how input queries, keys, and values are processed to produce context-aware outputs. This process underpins the model's ability to handle sophisticated language understanding and generation tasks efficiently." Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention FigureId="4.2" FigureContent="The image depicts the architecture of Multi-Head Attention, a key concept in transformer models used for natural language processing and other tasks. Let's break down the components and structure of this diagram step by step: Scaled Dot-Product Attention (Left Side): This is the fundamental building block of the attention mechanism. Input Elements: Queries (Q), Keys (K), and Values (V). The process involves calculating dot products of the query with all keys, scaling the result, and applying a softmax function to obtain the weights on the values. Multi-Head Attention (Right Side): This architecture allows the model to jointly attend to information from different representation subspaces at different positions. Components: Input Linear Layer: The input is linearly projected into multiple subspaces, producing multiple sets of Q, K, and V (typically, Q_1, Q_2, \dots, Q_h , K_1, K_2, \dots, K_h , and V_1, V_2, \dots, V_h). Attention Heads: Each set of Q, K, and V undergoes the Scaled Dot-Product Attention independently in parallel. This creates multiple different attention outputs. Concatenation Step: The output from each attention head is concatenated (Concat) to form a single matrix. Output Linear Layer: This concatenated matrix is linearly projected (via a linear layer) to the final output. Details: Notation (Bottom Labels): V, K, and Q represent the Values, Keys, and Queries fed into the attention mechanism. h indicates the number of parallel attention mechanisms, or heads, which is a key feature of multi-head attention. This allows for richer representations and more complex interactions within the model. Actionable Insights: Parallelism: Multi-head attention enhances the model's ability to focus on different positions for each head; hence, it captures various aspects of the data distribution. Efficiency: The scaled dot-product attention efficiently computes attention scores by leveraging matrix operations, making it suitable for training on large datasets."

search query : scaled
dot product attention

Documents68

Total storage2.18 MB

Vector index size413.93 KB

Max storage160 GB

Search explorer

Fields

CORS

Scoring profiles

Semantic configurations

Vector profiles

Query options

View

Scaled Dot Product Attention

Search

Results

1

{

2

"@odata.context": "https://aisearch-text2sql-adi.search.windows.net/indexes('rag-documents-index-old-test')/\$me

3

"@odata.count": 52,

4

"@search.answers": [],

5

"@search.nextPageParameters": {

6

"search": "Scaled Dot Product Attention"

Save

Discard

Refresh

Create demo app

Edit JSON

Delete

Documents47

Total storage1.7 MB

Vector index size286.74 KB

Max storage160 GB

Search explorer

Fields

CORS

Scoring profiles

Semantic configurations

Vector profiles

Query options

View

Scaled Dot Product Attention

Search

Results

7

/

8

"@search.score": 0.0333333350/180214,

9

"@search.rerankerScore": 3.202665328979492,

10

"@search.captions": [

11

"text": "To counteract this effect, we scale the dot products by a. #### 3.2.2 Multi-Head Attention

12

"highlights": "To counteract this effect, we scale the dot products by a. #### 3.2

Old Index

Search Results

Updated Index

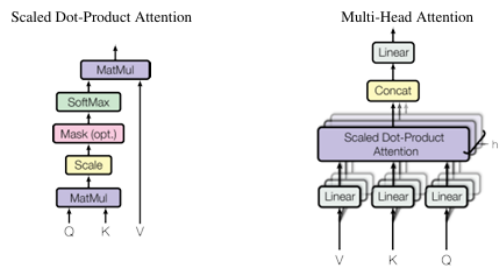


Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention consists of several attention layers running in parallel.

of the values, where the weight assigned to each value is computed by a compatibility function of the query with the corresponding key.

3.2.1 Scaled Dot-Product Attention

We call our particular attention "Scaled Dot-Product Attention" (Figure 2). The input consists of queries and keys of dimension d_k , and values of dimension d_v . We compute the dot products of the query with all keys, divide each by $\sqrt{d_k}$, and apply a softmax function to obtain the weights on the values.

In practice, we compute the attention function on a set of queries simultaneously, packed together into a matrix Q . The keys and values are also packed together into matrices K and V . We compute the matrix of outputs as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \tag{1}$$

The two most commonly used attention functions are additive attention [2], and dot-product (multiplicative) attention. Dot-product attention is identical to our algorithm, except for the scaling factor of $\frac{1}{\sqrt{d_k}}$. Additive attention computes the compatibility function using a feed-forward network with a single hidden layer. While the two are similar in theoretical complexity, dot-product attention is much faster and more space-efficient in practice, since it can be implemented using highly optimized matrix multiplication code.

While for small values of d_k the two mechanisms perform similarly, additive attention outperforms dot product attention without scaling for larger values of d_k [3]. We suspect that for large values of d_k , the dot products grow large in magnitude, pushing the softmax function into regions where it has extremely small gradients [4]. To counteract this effect, we scale the dot products by $\frac{1}{\sqrt{d_k}}$.

AutoFunction Calling

```
question = "Explain in detail scaled dot product attention"
await ask_question(question, history)
```

[12] ✓ 24.7s Python

```
[BEGIN ENTITY = &#x27;GET ALL CATEGORIES&#x27;]
  Name=&#x27;Get All Categories&#x27;
  Description=&#x27;This view provides a comprehensive list of
[END ENTITY = &#x27;GET ALL CATEGORIES&#x27;]

[BEGIN ENTITY = &#x27;PRODUCT AND DESCRIPTION&#x27;]
  Name=&#x27;Product and Description&#x27;
  Description=&#x27;This view provides detailed information ab
[END ENTITY = &#x27;PRODUCT AND DESCRIPTION&#x27;]

[BEGIN ENTITY = &#x27;PRODUCT MODEL CATALOG DESCRIPTION&#x27;]
  Name=&#x27;Product Model Catalog Description&#x27;
  Description=&#x27;This view provides detailed catalog inform
[END ENTITY = &#x27;PRODUCT MODEL CATALOG DESCRIPTION&#x27;]
```

```
question = "Comparison of phi-3-mini, phi-3.5-mini and phi-3.5-MoE on MMLU-Multilingual tasks"
await ask_question(question, history)
```

[14] ✓ 25.2s Python

```
... /tmp/ipykernel_61002/305922816.py:31: DeprecationWarning: Use the `FunctionChoiceBehavior` `Auto` class method instead.
      req_settings.function_call_behavior = FunctionCallBehavior.EnableFunctions(auto_invoke=True, filters=filter)
WARNING:semantic_kernel.connectors.ai.open_ai.prompt_execution_settings.open_ai_prompt_execution_settings:The `function_call_be
INFO:root:Question: Comparison of phi-3-mini, phi-3.5-mini and phi-3.5-MoE on MMLU-Multilingual tasks
INFO:semantic_kernel.functions.kernel_function:Function ChatBot-Chat invoking.
INFO:semantic_kernel.contents.chat_history:Could not parse prompt <message role="system">
As a senior analyst, your primary responsibility is to provide accurate, thorough answers to user queries. Use all available fu

- Always use the provided functions to obtain key information.
- If a function is required, you must use it to complement the answer.
- Use multiple functions in parallel to enhance the results.
- Always provide an answer; never leave it blank.

The response must meet the following requirements:

[RESPONSE OUTPUT REQUIREMENTS]

The answer MUST be in JSON format:
{
```


Challenges & Learnings

- Azure Document Intelligence unable to parse PPTX and XLSX files
- Moving from inbuilt skills for enrichment to custom skills
- Monitoring and debugging of each session for failures
- Moving to ML implementation for Search
- Metadata enrichment for index
- Page wise chunking retrieved better search results given the structure of majority documents , however semantic chunking can also be tried out



Text2SQL

Structured Data Search

Industry Solutions | Data & AI

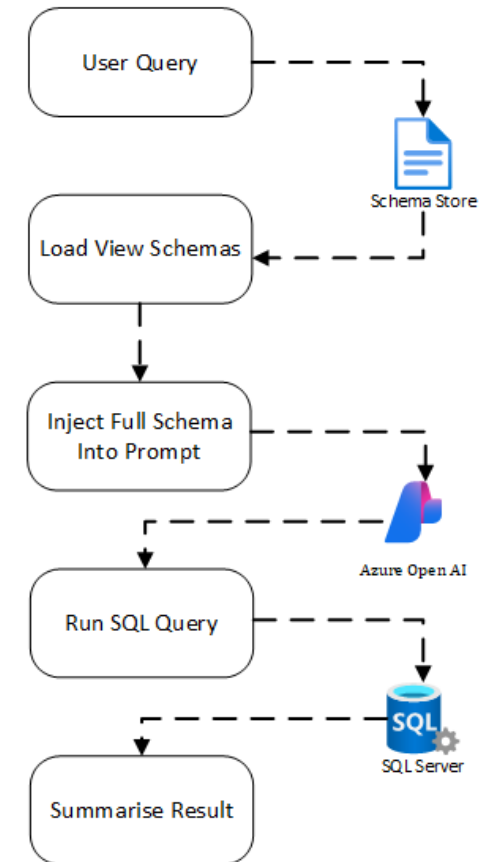
Why Text2SQL?

- Allows aggregations and calculations on data quickly
 - Without SQL integration, questions may need specifically pre-computed report
 - e.g. [What is our top performing sales per by quantity of units sold this month?](#) is answerable with a few simple SQL query if the correct views are exposed.
- Without Text2SQL, automatic report generation needs to be built.
- Pushing numerical calculations onto the source SQL engine ensures accuracy in the maths.
- Data can be updated real-time in the source database and be immediately accessible to the LLM.

Common Text2SQL Approach – Single Shot

- Limitations of single shot:
 - More tables / views significantly **increases the number of tokens** used within the prompt and the cost of inference.
 - Increasingly expensive operation
 - Excessive amounts of schema information can cause confusion with the LLM.
 - In our original use case, when **exceeding 5 complex tables / views**, we found that the LLM could **get confused** between which columns belonged to which entity and as such, would generate invalid SQL queries.

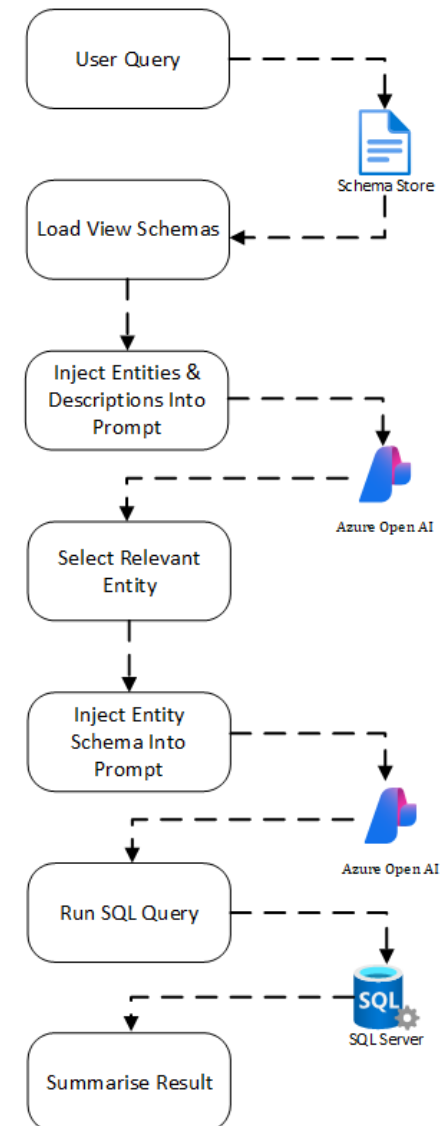
Common Text2SQL Approach



Our Initial Approach – Multi-Shot

- Break down the query generation into 2 steps:
 - Selection of the correct view(s)
 - Retrieval of the full schemas for the selected views
- Advantages:
 - Reduces token usage
 - Reduces confusion between schema attributes

Multi-Shot Text2SQL Approach



Data Dictionary & Metadata Gathering

- Maps out all the entities available in the SQL Source
 - Includes detailed descriptions for all columns available and the data types. Effectively “prompt engineering” per column
 - Avoids adding latency in fetching the schemas at runtime
- Quality of metadata **directly determines the quality of entity and column selection.**

```
{
  "EntityName": "Get All Categories",
  "Entity": "vGetAllCategories",
  "Description": "This view provides a comprehensive list of all product categories and their corresponding subcategories in the SalesLT schema of the AdventureWorksLT database. It is used to understand the hierarchical structure of product categories, facilitating product organization and categorization.",
  "Columns": [
    {
      "Definition": "A unique identifier for each product category. This ID is used to reference specific categories.",
      "Name": "ProductCategoryID",
      "Type": "INT"
    },
    {
      "Definition": "The name of the parent product category. This represents the top-level category under which subcategories are grouped.",
      "Name": "ParentProductCategoryName",
      "Type": "NVARCHAR(50)"
    },
    {
      "Definition": "The name of the product category. This can refer to either a top-level category or a subcategory, depending on the context.",
      "Name": "ProductCategoryName",
      "Type": "NVARCHAR(50)"
    }
  ]
}
```

Automated Metadata Gathering

- Metadata collection is a time-consuming exercise.
- **Automatic metadata collection** is used with a combination of SQL queries and prompting with an LLM.

```
"Entity": "SalesLT.vGetAllCategories",
>Description": "The SalesLT.vGetAllCategories entity provides a comprehensive view of all product categories and their hierarchical relationships within a sales database. It includes information on product category names, their parent category names, and unique identifiers for each category. This entity is useful for answering questions related to the organization and structure of product categories, such as identifying parent-child category relationships or retrieving specific product category details based on their IDs.",
>EntityName": "Product Categories Overview",
>Columns": [
>{
>Name": "ParentProductCategoryName",
>Type": "nvarchar",
>Definition": "The column ParentProductCategoryName in the SalesLT.vGetAllCategories entity contains the names of top-level product categories. The values represent broad categories that group various related products together. Examples of these categories include Components, Clothing, Bikes, and Accessories. This column is useful for identifying and organizing products under their respective high-level classifications.",
>AllowedValues": null,
>SampleValues": [
>{
>Components",
>Clothing",
>Bikes",
>Accessories"
>}
]},
>{
>Name": "ProductCategoryName",
>Type": "nvarchar",
>Definition": "The ProductCategoryName column in the SalesLT.vGetAllCategories entity contains the names of different product categories. The categories range from various bicycle components and accessories to cycling apparel. These names are typically descriptive of the type of product they represent, such as \"Pumps\" or \"Touring Frames.\" This column helps in organizing products into specific categories for easier identification and filtering.",
>AllowedValues": null,
>SampleValues": [
>{
>Pumps",
>Bottles and Cages",
>Locks",
>Touring Frames",
>Caps"
>}
]}
],
}
```

Prompt Injection

- Data dictionary transformed at runtime into formatted prompt.

```
[BEGIN ENTITIES LIST]
  [BEGIN ENTITY = 'SALES ORDER DETAIL']
    Name='Sales Order Detail'
    Description='This table stores detailed information about sales order tickets, including the order details, customer
    information, order status, and timestamps. It is used to manage and track sales orders throughout the order lifecycle, from
    creation to fulfillment.'
  [END ENTITY = 'SALES ORDER DETAIL']

  [BEGIN ENTITY = 'SALES ORDER HEADER']
    Name='Sales Order Header'
    Description='This table contains high-level information about sales orders, including order dates, customer details, shipping
    information, and order status. It is used to manage and track sales orders from initiation to fulfillment.'
  [END ENTITY = 'SALES ORDER HEADER']

  [BEGIN ENTITY = 'ADDRESS']
    Name='Address'
    Description='This table stores address information for customers, including street addresses, city, state, postal code, and
    country/region. It is used to maintain contact and shipping information for orders, as well as to manage customer locations.'
  [END ENTITY = 'ADDRESS']
[END ENTITIES LIST]
```

- **GetEntitySchema()** method is exposed to the LLM.
 - LLM selects and calls **GetEntitySchema()** via Auto Function calling to retrieve the full schema for the entities it believes will answer the question.
 - Reduces the passing of all schemas to the LLM, to only those that match the user's question.

Text2SQL Examples

What is the top performing quantity in terms of sales numbers?

```
"answer": "The top-performing product by quantity of units sold is the **Classic Vest, S** from the **Classic Vest** product model, with a total of 87 units sold [1][2].",
"sources": [
  {
    "title": "Sales Order Detail",
    "chunk": "| ProductID | TotalUnitsSold |\n|-----|-----|\n| 864      | 87      |\n",
    "reference": "SELECT TOP 1 ProductID, SUM(OrderQty) AS TotalUnitsSold FROM SalesLT.SalesOrderDetail GROUP BY ProductID ORDER BY TotalUnitsSold DESC;"
  },
  {
    "title": "Product and Description",
    "chunk": "| Name          | ProductModel |\n|-----|-----|\n| Classic Vest, S| Classic Vest |\n",
    "reference": "SELECT Name, ProductModel FROM SalesLT.vProductAndDescription WHERE ProductID = 864;"
  }
]
```

The top-performing product by quantity of units sold is the **Classic Vest, S** from the **Classic Vest** product model, with a total of 87 units sold [1][2].

Rendered Sources

ProductID	TotalUnitsSold
864	87

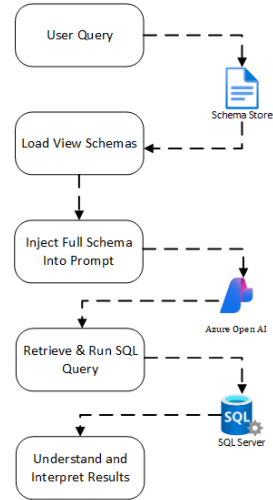
Name	ProductModel
Classic Vest, S	Classic Vest

Further Improvements

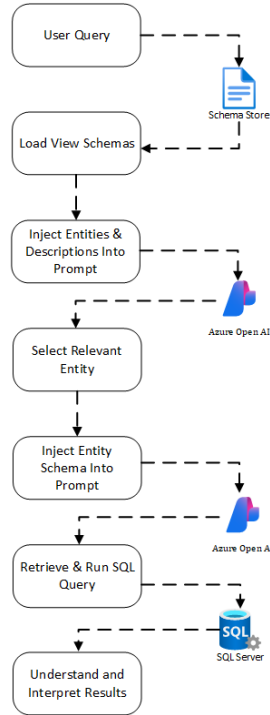
- Offload selection of entities from the LLM.
 - Reduce token usage
 - Speed up selection
 - Improve reliability by using deterministic approach
- Cache the question to query resolution.
 - Speed up resolution for commonly asked questions.
- Pre-fetch database results for commonly asked questions
 - Can we [skip query generation completely](#)?

Text2SQL Further Iterations

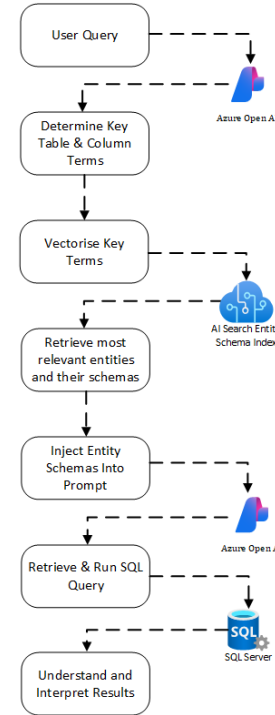
Common Text2SQL Approach
(Iteration 1)



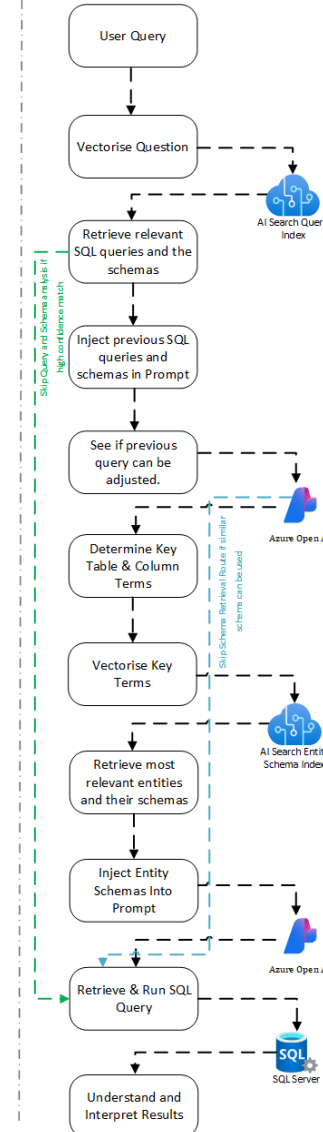
Prompt Based Multi-Shot Text2SQL Approach
(Iteration 2)



Vector Based Multi-Shot Text2SQL
(Iteration 3)













Vector Based Multi-Shot Text2SQL With Query Cache
(Iteration 4)



Schema Index Setup



- Transforms the data dictionary into a index.
- Vectorises the description property.
- Setups column names as "keywords" in semantic config.

	Field name	Type	Retrievable	Filterable	Sortable	Facetable	Searchable	Analyzer	Dimensions
			<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>		
	Entity	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	Keywo... 	
	EntityName	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	
	Description	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	
	DescriptionEmbedding	SingleCollection	<input type="checkbox"/>				<input checked="" type="checkbox"/>		1536
	Columns	ComplexTypeCollection							
	Name	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	
	Definition	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	
	Type	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	
	AllowedValues	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input type="checkbox"/>		
	SampleValues	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input type="checkbox"/>		
	ColumnNames	StringCollection	<input type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... 	

- Search term of the original question, automatically resolves to the most relevant entities.

Query Cache Index Setup

- Stores the user's question and the queries / schemas it resolves to.
- Vectorises the question property.

Field name	Type	Retrievable	Filterable	Sortable	Facetable	Searchable	Analyzer	Dimensions
		<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>		
 Id	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>		
Question	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	Keywo... ▾	
 QuestionEmbedding	SingleCollection	<input type="checkbox"/>				<input checked="" type="checkbox"/>		1536
Query	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... ▾	
⌵ Schemas	ComplexTypeCollection							
Entity	String	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... ▾	
⌵ Columns	ComplexTypeCollection							
Name	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... ▾	
Definition	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... ▾	
Type	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Standa... ▾	
AllowedValues	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input type="checkbox"/>		
SampleValues	String	<input checked="" type="checkbox"/>	<input type="checkbox"/>		<input type="checkbox"/>	<input type="checkbox"/>		
DateLastModified	DateTimeOffset	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>			

- Ensures consistency that the same question, will always resolve to the same schemas.

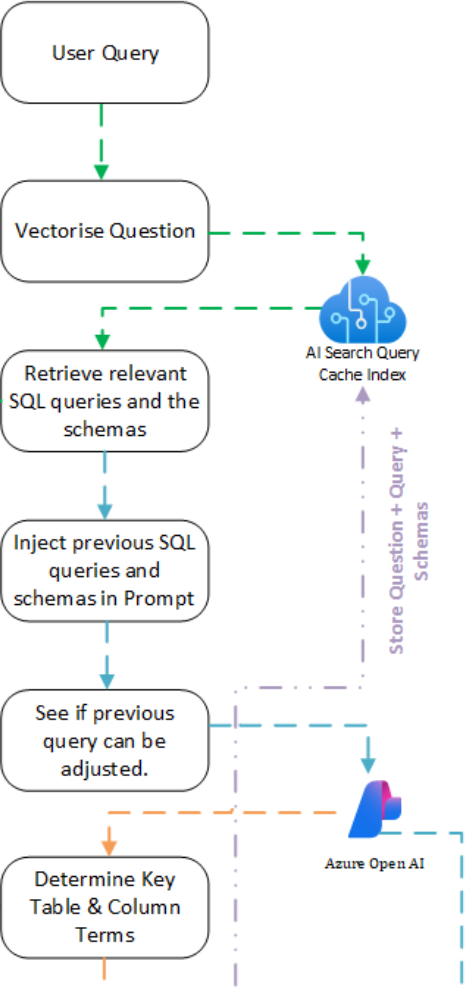
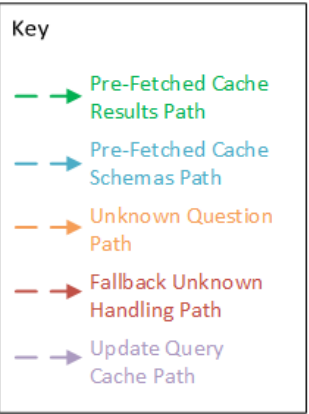
Query Cache Index Setup

```
{
  "@search.score": 0.8020883,
  "Id": "R2l2ZSBtZSB0aGUgdG90YWwgbnVtVmVyIG9mIG9yZGVycyBpbIAyMDA4Pw==",
  "Question": "Give me the total number of orders in 2008?",
  "Query": "SELECT COUNT(SalesOrderID) as TotalOrders FROM SalesLT.SalesOrderHeader WHERE YEAR(OrderDate) = 2008;",
  "DateLastModified": "2024-09-20T19:33:47.461Z",
  "Schemas": [
    {
      "Entity": "SalesOrderHeader",
      "Columns": [
        {
          "Name": "SalesOrderID",
          "Definition": "A unique identifier for each sales order. This ID is auto-generated and serves as the primary key for the SalesOrderHeader table.",
          "Type": "INT",
          "AllowedValues": null,
          "SampleValues": null
        },
        {
          "Name": "OrderDate",
          "Definition": "The date and time when the sales order was created. This field is used to track when the order was initiated.",
          "Type": "DATETIME",
          "AllowedValues": null,
          "SampleValues": null
        },
        {
          "Name": "DueDate",
          "Definition": "The date by which the order is expected to be fulfilled or delivered. It helps in managing delivery timelines.",
          "Type": "DATETIME",
          "AllowedValues": null,
          "SampleValues": null
        },
        {
          "Name": "ShipDate",
          "Definition": "The date when the order was shipped to the customer. This is used for tracking shipping and fulfillment status.",
          "Type": "DATETIME",
          "AllowedValues": null,
          "SampleValues": null
        },
        {
          "Name": "Status",
          "Definition": "The current status of the order, represented as a numeric code (e.g., 1 for In Progress, 2 for Completed, 3 for Canceled).",
          "Type": "TINYINT",
          "AllowedValues": "[1,2,3]",
          "SampleValues": null
        },
        {
          "Name": "OnlineOrderFlag",
          "Definition": "Indicates whether the order was placed online.",
          "Type": "BIT",
          "AllowedValues": "[\"True\", \"False\"]",
          "SampleValues": null
        }
      ]
    }
  ]
}
```

Query Caching

Full Flow for Vector Based Multi-Shot Text2SQL With Query Cache

(Iteration 4)

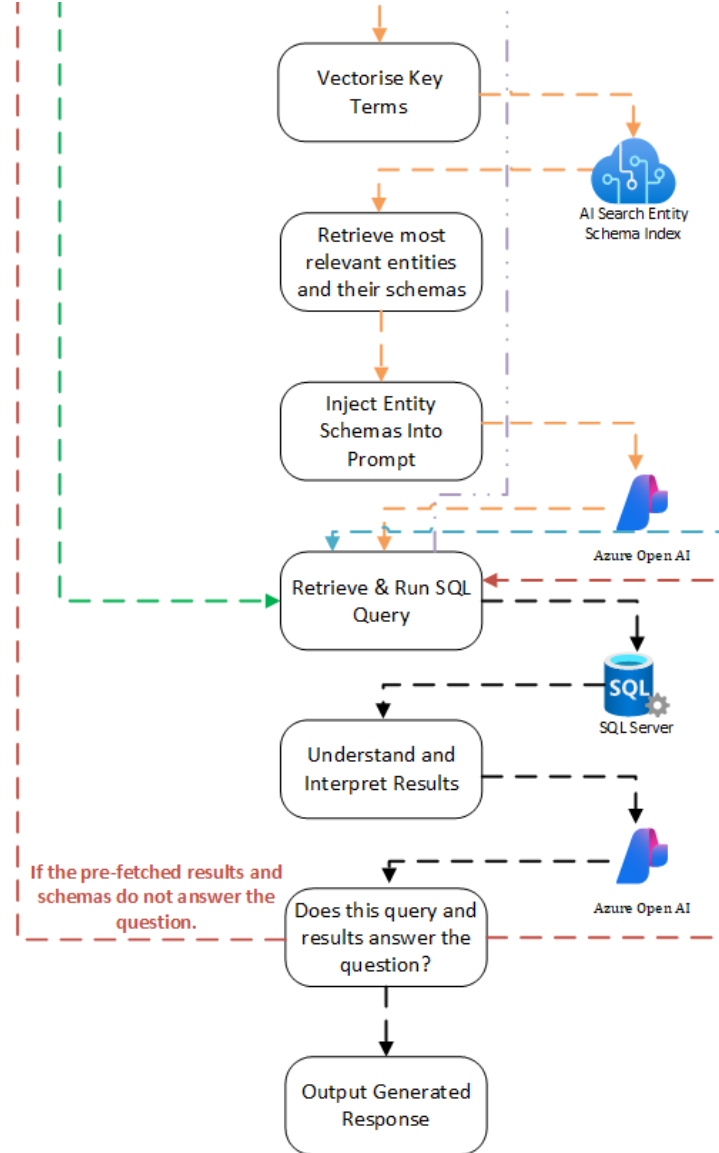



Skip Query and Schema analysis if high confidence match

Store Question + Query + Schemas

Full Flow for Vector Based Multi-Shot Text2SQL With Query Cache

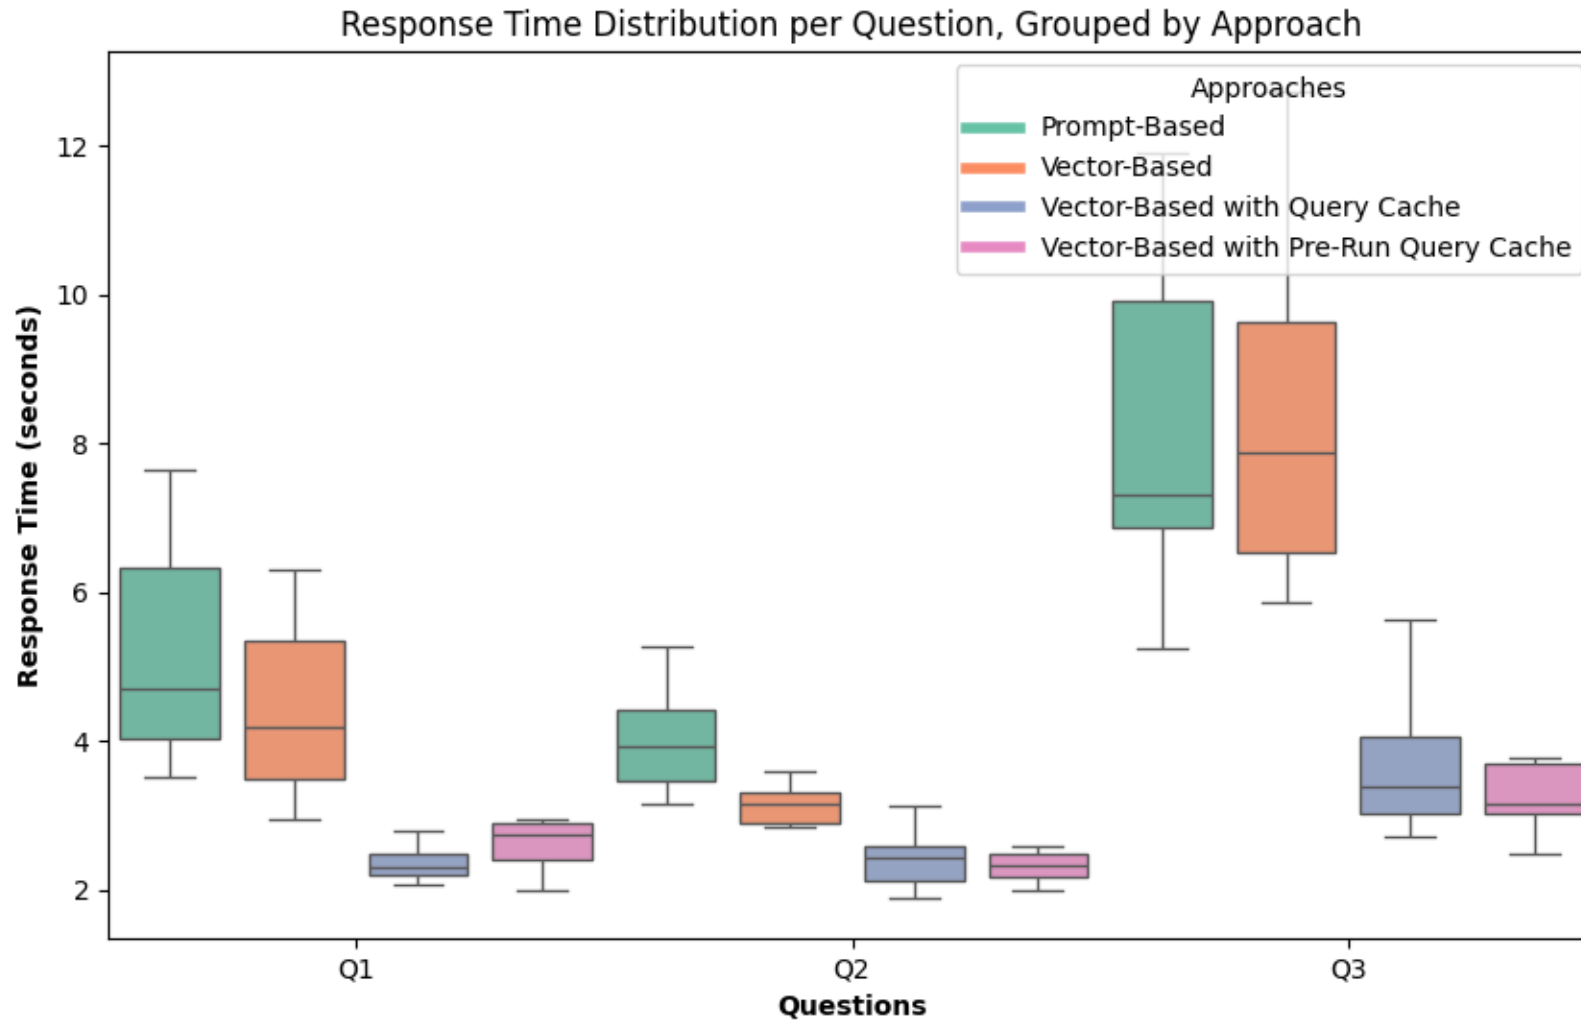
(Iteration 4)



Skip Schema Retrieval Route if similar schema can be used

If the pre-fetched results do not answer the question but the query can be adjusted based on the schemas

Timing Comparison



Text2SQL Challenges

- Selecting the best entity / column in cases where multiple columns may be a good match.
 - Similarly named columns may seem relevant for a given question.
- Text representation
 - Different columns are formatted in different ways, achieving accurate SQL statements requires knowledge of the source system representation.
- Performance
 - Initial approach was slow and struggled to scale with multiple tables.

Text2SQL Learnings & Performance Tips

- Pre-assemble views to avoid the LLM having to make complex joins between multiple tables
- Spend time providing detailed names and descriptions in the metadata for all entities and columns e.g.
 - If a column contains a value in each currency, give the currency information in the metadata.
 - Clearly state in the description what sorts of questions a given view / table can provide answers for.
- Use common codes for columns that need filtering e.g.
 - A country can have multiple text representations e.g. United Kingdom or UK.
 - Use ISO codes for countries, instead of text descriptions to increase the likelihood of correct and valid SQL queries.

IP Available

Full Python Code available as part of the Data Science Toolkit:

- Function App for Document Extraction With Azure Document Intelligence
- Text2SQL Plugin for both Prompt & Vector Based Approaches (including cache tweaks)
- Deployment scripts for index / indexer setup

aka.ms/text2sql-adi



Thank you

Industry Solutions | Data & AI